


Interaction memory supplements a stable visual-language state

PROPOSED METHOD • all memory evidence precedes the decision it conditions



Longer persistence is explicit: verified retries or a separate source bank; no unrestricted lifetime archive.

Stage 1B: learn to read past interactions

Past context + current state + target action
→ predict the subsequent joint / visual outcome
No target outcome in the context

Stage 2: evaluate control value

Freeze backbone and pretrained context first; train actor / critics
Store immutable decision-time snapshots in replay
Frozen-weight evaluation: retain vs. clear vs. shuffled memory