


(a) Stage 1B: learn a context that predicts a later response



(b) Stage 2: fixed representation, small actor-critic updates

